

Shihab Shahriar Khan

Personal: shihab-shahriar.github.io

GitHub: /shihab-shahriar

Google Scholar: Profile Link

Email: khanmd@msu.edu

EXPERIENCE

Lam Research

Software Development Intern

Austin, TX

May 2025 – Aug. 2025

- Introduced shared memory parallelism to Lam's Vizglow software (250k+ LOC C++ plasma sim): achieved **2-4x speedup** over existing MPI-only code. Built thread-safe sparse data structures, parallelized mesh traversals, optimized memory allocations, std library calls overhead etc. *This work has been accepted into Vizglow's next-release roadmap.*
- Ported the most time-consuming sparse linear solver task to GPU: removed a critical library dependency by isolating API calls and generalizing the solver interface; minimized CPU-GPU comms, overlapping with computation etc. From 3x-20x speedup compared to CPU baseline (Currently in exploratory phase).

Los Alamos National Laboratory

Graduate Research Intern

Los Alamos, NM

June 2024 – Aug. 2024

- Completed the development of an open-source Hydrodynamics and Radiative Diffusion Solver (HARD), achieving performance comparable to state-of-the-art solutions. ([Publication](#))
- I led the parallelization and GPU acceleration of the program, working as a member of a diverse team of astrophysicists and mathematicians, utilizing C++, Kokkos, MPI, CUDA, and Nsight profilers.

Michigan State University

PhD Candidate

East Lansing, MI

Mar. 2022 – Present

- Developing a performance-portable simulation framework for particulate flows, designed for seamless scaling across multi-node systems and on-node with GPU acceleration.
- Exploring how machine learning techniques can accelerate the compute-intensive evaluation of N-body hydrodynamic interactions. Tech stack includes C++, CUDA, OpenMP/MPI, Kokkos, Python, Pytorch.

PUBLICATIONS

1. **Khan, Shihab S**, Bryce Palmer, Christopher Edelmaier, and Hasan M Aktulga. “OpenRAND: A performance portable, reproducible random number generation library for parallel computations.” *SoftwareX* 27 (2024): 101773.
2. Loiseau, Julien, Hyun Lim, Andrés Yagüe López, Mammadbaghir Baghiazade, **Shihab Shahriar Khan**, Yoonsoo Kim, Sudarshan Neopane, Alexander Strack, Farhana Taiyebah, and Benjamin K. Bergen. “HARD: A Performance Portable Radiation Hydrodynamics Code based on FleCSI Framework.” *SoftwareX* Vol 32 (2025): 102441.
3. **Shihab Shahriar Khan**, Nishat Tasnim Niloy, Md. Aquib Azmain and Ahmedul Kabir. “Impact of Label Noise and Efficacy of Noise Filters in Software Defect Prediction”. *International Conference on Software Engineering and Knowledge Engineering (SEKE)* KSIR Virtual Conference Center, USA, 2020.

EDUCATION

Michigan State University

- Ph.D. in Computer Science & Engineering
 - Supervised by: Dr Hasan M Aktulga
 - CGPA: 3.92/4.00

MI, USA

Aug. 2021 – Present

Institute of Information Technology, University of Dhaka

- M.S. in Software Engineering
 - Thesis Title: “Noise Robust Classification Using Instance Level Analysis”. CGPA: 3.83/4.00

Dhaka, Bangladesh

Jan. 2019 – Dec. 2020

Institute of Information Technology, University of Dhaka

- B.S. in Software Engineering
 - CGPA: 3.68/4.00

Dhaka, Bangladesh

Jan. 2015 – Dec. 2018

TECHNICAL SKILLS

High Performance Computing

Tech Stack: C/C++, CUDA, Kokkos, OpenMP/MPI, x86/GPU architecture, NSight Systems/Compute, Programming Models: SIMT/SIMD/Multithreaded/Distributed Parallelism

Machine Learning and Data Science

Python, Pytorch, Triton, Numpy/Pandas, R, Scipy, Scikit-learn, SQL, Tensorflow

Common Skills and Tools

Advanced data structures, Technical writing, Git, Linux, CI/CD, CMake/Make, Haskell, Java

PROJECTS

OpenRAND ([Link](#))

Reproducible random number generation for parallel computations (Lead Developer) Sept. 2023 – Present

- An open-source, statistically robust C++17 library aimed at facilitating reproducible scientific research through the generation of reproducible, parallel random number streams
- Performance-portable, designed to work seamlessly across various software and hardware platforms (including GPUs.) and often faster than native libraries like libstdc++, Nvidia’s Curand or rocRAND.
- Merged into HOOMD-blue, slated for inclusion in Mundy (from Flatiron Institute).

HARD ([Link](#))

Hydrodynamics And Radiative Diffusion Solver based on FleCSI framework June 2024 – Jan 2025

- The first publicly available code coupling radiation and hydrodynamics that is highly scalable and portable to heterogeneous HPC architectures.
- GPU acceleration led to around 15 times performance improvement over a 64-core intel CPU on a A100 GPU, resulting in more than twice cell updates per second in single node configuration over best known literature result.

Scikit-clean ([Link](#))

A python ML library for classification in the presence of label noise. Dec. 2019 – Dec. 2021

- Tools to simulate artificial noise, create complex pipelines and evaluate them
- scikit-learn API compatible- all scikit-learn’s building blocks can be seamlessly integrated into workflow
- Equipped with units tests, extensive documentation, CI pipeline and available in pypi.

DeepPaint ([Link](#))

Deep learning based computer vision tool to automatically colorize and stylize paintings Jul. 2018 – Nov. 2018

- Colorizes a sketch using color hints. Stylizes a colored painting along the style of any given painting.
- Implemented using python, pytorch, pyqt.

ACTIVITIES

Open Source

Several code-level pull requests accepted in projects like Kokkos, Spack, Imbalance-learn, and MLAlgorithms. OpenRAND is currently part of HOOMD-blue, slated for inclusion in Mundy (from Flatiron Institute), possibly Stan.

Technical Writing Samples

- Article: “CUDA vs ROCm: A Case Study”
- Article in Towards Data Science titled “An Introduction to Classification Using Mislabeled Data”.
- StackOverflow: Old Stackoverflow answers, mostly in Python, Scikit-learn, Pytorch and Algorithm tags.